

# decision trees

## basics

- test & prediction
- input & output:

- We will denote the input space by  $X$ ; points in  $X$  are vectors  $\mathbf{x} = (x_1, x_2, \dots, x_n)$
- Any  $\mathbf{x}$  is mapped to exactly one value  $y$  for  $Y$
- Hence, the tree represents a function from  $X$  to  $Y$

- representation power

can every imaginable boolean function be represented?

yes

- continuous input attributes

cannot make a different child node for each possible value!

solution: comparative tests

- classification tree & regression tree

- why? (advantages)

efficient

good predictive accuracy

interpretable

## the basic learning algorithm

### task 1: smallest

Find the smallest tree  $T$  such that  $\forall (\mathbf{x}, f(\mathbf{x})) \in D : T(\mathbf{x}) = f(\mathbf{x})$   
(= smallest tree *consistent* with the data)

- smallest tree = simplest explanation, provides insight in the data
- not practical,  $D$  is only a sample from some larger distribution (lack of generalization)

## task 2: minimal risk (loss)

- loss function
- expected loss

- The **risk**  $R$  of  $T$ , relative to  $f$ , is  $\mathbf{E}_{\mathbf{x} \sim \mathcal{D}}[(\ell(T(\mathbf{x}), f(\mathbf{x}))]$   
(the expected value of  $\ell(T(\mathbf{x}), f(\mathbf{x}))$ , with  $\mathbf{x}$  drawn from  $\mathcal{D}$ )

using only the data in  $D$ , we want both aspects to be achieved (is small & generalizes well)

## learning algorithm

- hardness of learning decision trees: NP-hard
- the basic principle: TDIDT, recursive partitioning

- Start with the full data set  $D$
- Find a test such that examples in  $D$  with the same outcome for the test tend to have the same value of  $Y$
- Split  $D$  into subsets, one for each outcome of that test
- Repeat this procedure on each subset that is not yet sufficiently “pure” (meaning, not all elements have the same  $Y$ )
- Keep repeating until no further splits possible

- two important questions:
  - how to choose the "best" test
  - when to stop splitting nodes

## choosing a test: classification

- information entropy

- Given a set of values  $c_1, c_2, \dots, c_k$  with respective probabilities  $p_1, p_2, \dots, p_k$ , an encoding exists that uses, on average,  $e$  bits for representing a randomly drawn value, where

$$e = - \sum_{i=1}^k p_i \log_2(p_i)$$

$e$  reflects the minimal number of bits that you will need (on average) to encode one value

#### ■ class entropy

- The **class entropy** of a set  $S$  of objects  $(\mathbf{x}, y)$ , where  $y$  can be any of  $k$  classes  $c_i$ , is defined as

$$CE(S) = - \sum_{i=1}^k p_i \log_2(p_i) \quad \text{with } p_i = \frac{|\{(x, y) \in S \mid y = c_i\}|}{|S|}$$

(proportion of elements in  $S$  with class  $c_i$ )

high entropy = (high uncertainty)

"many possibilities, all equally likely"

low entropy =

"few possibilities" /

"many possibilities but most are highly unlikely"

#### ■ information gain

expected reduction of entropy by obtaining the answer to a question

- In the case of classification trees: expected reduction of class entropy:

$$IG(S, t) = CE(S) - \mathbf{E}[CE(S_i)] = CE(S) - \sum_{i=1}^o \frac{|S_i|}{|S|} CE(S_i)$$

with  $t$  a test,  $o$  the number of possible outcomes of  $t$ , and  $S_i$  the subset of  $S$  for which the  $i$ 'th outcome was obtained

this measurement is for: given **a dataset** and **a test**!

## choosing a test: regression

goal: examples in one subset have similar Y values

use variance reduction instead of information gain:

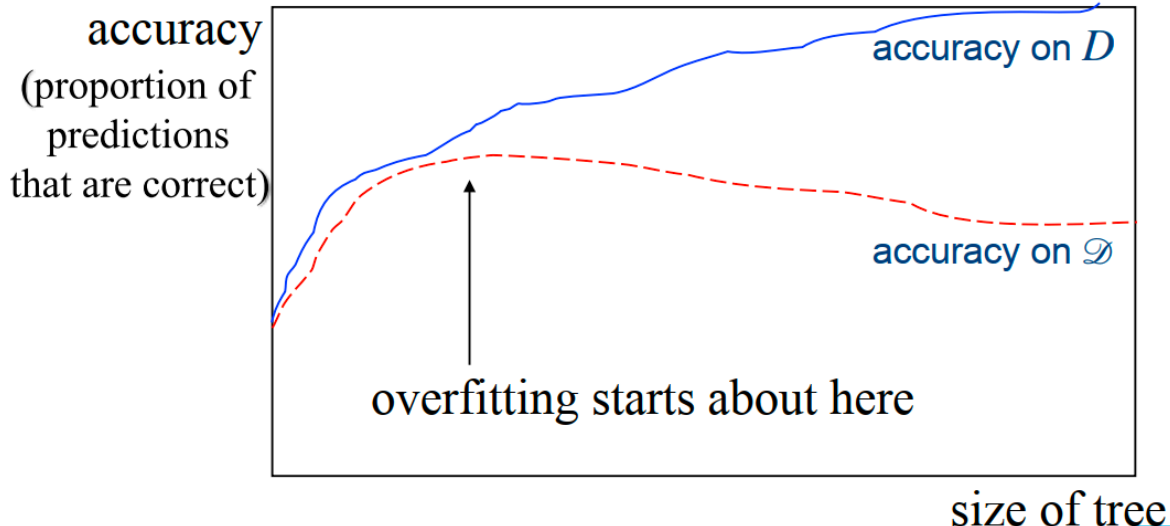
$$Var(S) = \frac{\sum_{(x,y) \in S} (y - \bar{y})^2}{|S| - 1} \quad \text{with } \bar{y} = \frac{\sum_{(x,y) \in S} y}{|S|}$$

$$VR(S, t) = Var(S) - \sum_{i=1}^o \frac{|S_i|}{|S|} Var(S_i)$$

## stopping criteria

until all instances in a subset have the same Y value

-> useful for classification, but overfitting for regression!



to avoid overfitting,

### 1. cautious splitting

use a **validation set** to guess if the model is going to be overfitted

but what if the guess is wrong? this is not always reliable

## 2. post-pruning

grow the tree to full size, then cut away branches that don't contribute to getting better predictions

顺序：叶节点到高层节点


## a generic TDIDT algorithm

```
function TDIDT( $E$ : set of examples) returns tree;
     $T' = \text{grow\_tree}(E)$ ;
     $T = \text{prune\_tree}(T')$ ;
    return  $T$ ;

function grow_tree( $E$ : set of examples) returns tree;
     $T = \text{generate\_tests}(E)$ ;
     $t = \text{best\_test}(T, E)$ ;
     $P = \{E_1, E_2, \dots, E_k\}$  with  $E_i = \{x \in E \mid t(x) = v_i\}$ 
    (call  $t$ 's outcomes  $v_1 \dots v_k$ )
    ( $P$  = partition induced on  $E$  by  $t$ )
    if stop_criterion( $E, P$ )
    then return leaf(info( $E$ ))
    else
        for all  $E_i$  in  $P$ :  $T_i := \text{grow\_tree}(E_i)$ ;
        return node( $t, \{(v_1, T_1), (v_2, T_2), \dots (v_k, T_k)\}$ );
```

### ▪ generate\_tests

variants: one subtree per set of values, **oblique trees**





in practice, non-oblique trees are more common

- best\_test

**impurity** 不纯度 should be the reduction goal of building a tree

SEMINAL conclusion: good impurity measures are strictly concave





Accuracy will only improve if at least one child node has a majority class different from that of its parent

但这通常too restrictive:

一些good tests虽然不是严格的非零增益，却还是能降低数据的不纯度

for classification trees

information gain:

Gini impurity reduction:

$$Gini(S) = 1 - \sum_{i=1}^k p_i^2$$

gain ratio

= IG / class entropy:

- For a test  $t$  that splits  $S$  into  $n$  subsets  $S_i$ :

$$SI(S, t) = - \sum_{i=1}^n \frac{|S_i|}{|S|} \log_2 \frac{|S_i|}{|S|} \quad (\text{Note: this is the "t-entropy" in } S; \text{ cf. class entropy})$$

- The Gain Ratio is:  $GR(S, t) = IG(S, t) / SI(S, t)$

for regression trees

how to decide the  $c$  in " $x < c$ "?

typically, just try all values, complexity -- in next section

- stop\_criterion

post-pruning is always preferred!

still, early drop out

- info (in leaf nodes)

for classification trees

most frequent class in this leaf / class distribution / all training examples relevant for that leaf

for regression trees

mean / median ...

- prune\_tree

order matters

## complexity

tree construction 的主要计算量：

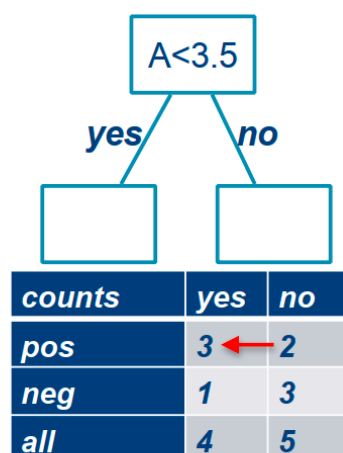
split每个节点的时候，为每个可能的test，partition数据集并计算在此partition上的性质

对于取值离散的classification任务，这个计算是简单的；

对于取值连续的回归任务，尝试all values：

但实际上复杂度并不高——首先根据分类标准A将样本排序，移动threshold即可

A	B	C	Class
2			neg
3			pos
3			pos
3			pos
4			neg
5			pos
6			pos
7			neg
8			neg



$|S_{\text{yes}}|=1$ ,  $|S_{\text{no}}|=8$   
 no:  $p_{\text{pos}}=2/5$ ,  $p_{\text{neg}}=3/5$   
 yes: ...

KU LEUVEN

- 预测是非常快的（理想情况下，如果树的高度不高）；而对于训练：
- 分裂一个节点的复杂度：



节点上的样本数 $n$ （不是总样本数 $N$ ），属性数量 $m$

则复杂度( $mn$ )（对每个属性计算，每次计算遍历 $n$ 个样本）

- 分裂多个节点的复杂度：

属性数量不变， $m * n = m * (n_1 + n_2) = m * n_1 + m * n_2$

所以无论怎么分裂，树的每一层总工作量都是几乎不变的

- 总复杂度

对于平衡的树，总复杂度就是 $O(m * N * \log N)$

对于不太平衡的树，高度甚至可能是线性的，总复杂度 $O(m * N^2)$

好的分裂heuristic，比如IG，倾向于构建出更平衡的树

## handling missing values

遇到缺省值

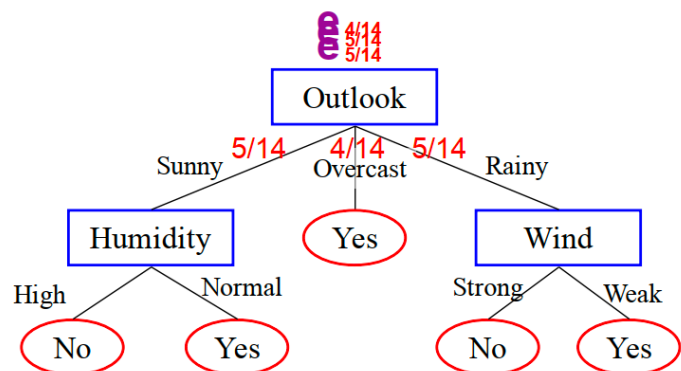
计算test quality的时候，直接跳过该样本即可

partition数据集的时候，两种方案：

1. 插值

2. distribution：一个样本在即将split的属性上有缺省值，将其加权重分配到不同分支上：

	Outlook	Hum.	Wind	Play
1	Sunny	High	Weak	No
2	sunny	high	strong	No
3	overcast	High	weak	yes
4	rain	High	Weak	yes
5	rain	normal	weak	yes
6	rain	normal	strong	no
7	overcast	normal	strong	yes
8	sunny	high	weak	no
9	Sunny	normal	weak	yes
10	rain	normal	weak	yes
11	Sunny	normal	strong	yes
12	overcast	high	Strong	yes
13	overcast	normal	weak	yes
14	rain	high	strong	No



e: 10/14 no, 4/14 yes => guess no

How to classify e : [15, ?, high, strong] ?

如果是预测阶段，直接加权投票得到预测结果；

如果是训练阶段，以0.3的权重分配到分支1算“0.3个”样本来计算基尼系数等；

## model trees

### each leaf contains a linear model

方差reduction假设每个叶节点是一个常数，不适用于包含一个线性模型的叶节点；

RETIS将其换成：对每个subset计算线性回归，然后计算平均预测残差的减少；但当属性数量非常多的时候效率很低；

mauve对此优化：计算 单变量线性回归 后的平均残差；

## multi-target trees

to predict multiple labels, 3 main approaches:

1. binary relevance: learn a binary tree (yes / no) for each label
2. label powersets: build one tree that predicts the combined label
3. vector encoding: variance of vectors in `best_test` , mean vector in `info`

hierarchical multilabel classification (HMC problem)

a label can only occur in a set if all its ancestors also occur