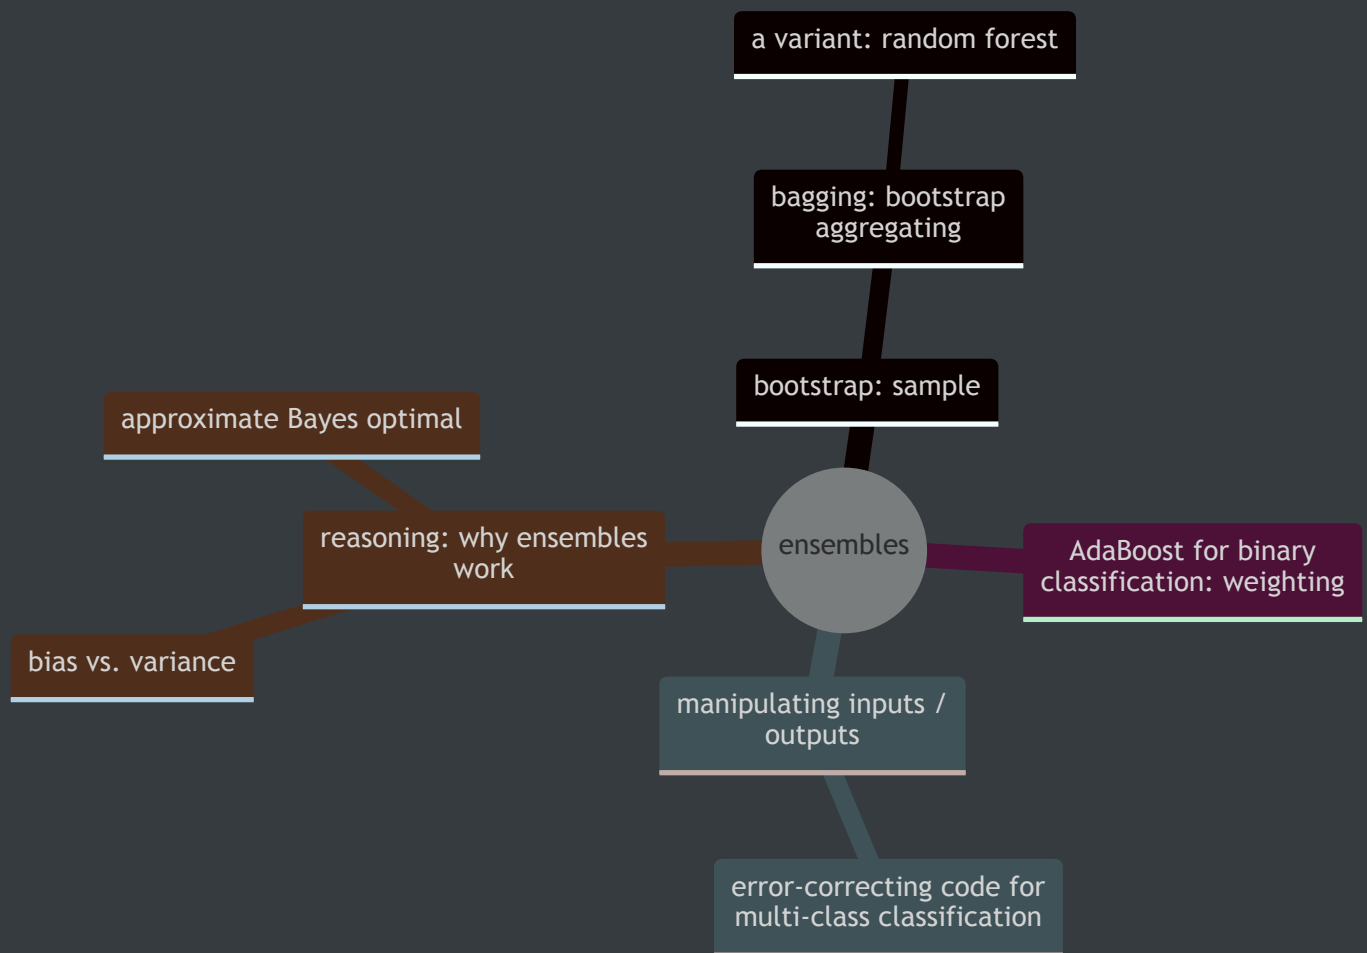


ensemble methods

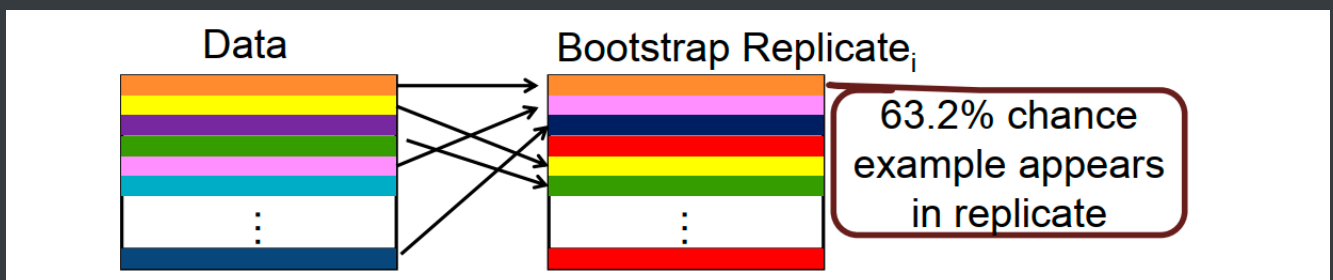
motivation & overview

can learning more than one model improve performance?



bagging & random forests

- bootstrap



Claim: 63.2% chance example x is in a bootstrap replicate S'

Proof:

$$P(x = S'_i) = \frac{1}{n}$$

x sampled for i^{th} position in S'

$$P(x \neq S'_i) = \left(1 - \frac{1}{n}\right)$$

x NOT sampled for i^{th} position in S'

$$P(x \in S') = \left(1 - \left(1 - \frac{1}{n}\right)^n\right)$$

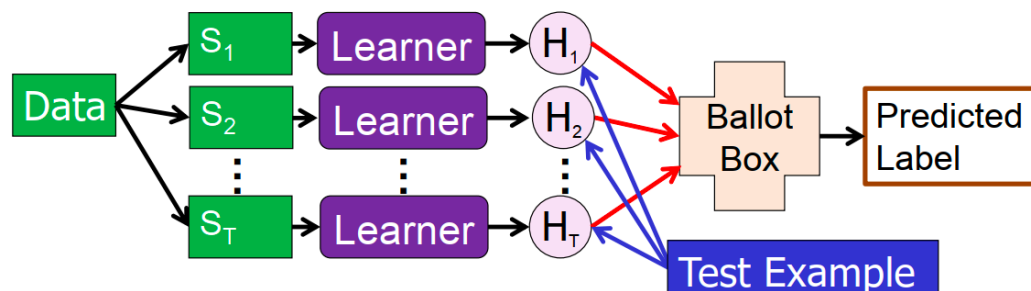
x NEVER sampled for S'

$$P(x \in S') \approx \left(1 - \frac{1}{e}\right) \approx 0.623$$

- Bagging: **B**ootstrap **A**ggregating

Given: Data set S , integer T

- For $t = 1, \dots, T$
 - ▣ S_t = bootstrap replicate of S
 - ▣ h_t = Apply learning algorithm to S_t
- Classify test instance using unweighted vote



c) jesse davis

- Bagging works best for unstable learners

unstable learners:

minor variations in training data \rightarrow major changes in classifier output

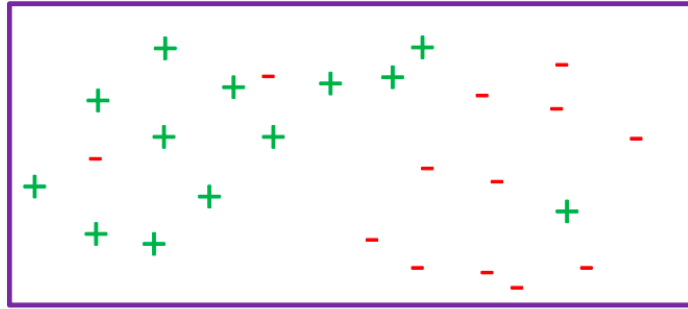
e.g.

unstable: decision tree, neural network, rule learning

stable: linear regression, nearest neighbor, linear threshold algorithms etc.

- intuition: why does bagging work? suppose for a noisy training data

- Suppose your data looks like this:



noisy samples might not appear in the samples used for training

(although kNN is stable learning)

- for Empirical Confidence Bounds

- Repeat 1000 (or 10,000) times:
 - ▣ Draw bootstrap sample
 - ▣ Repeat entire cross-validation process
 - ▣ Lower (upper) bound is result (e.g., accuracy) such that 2.5% of runs yield lower (higher)

then confidence level = 95%

- Random Forests

A variant of BAGGING

Given: Data set S , integer T

Algorithm

Repeat T times

1. S_t = bootstrap replicate of S
2. Build d-tree on S_t , **but** in each recursive call
 - A. Choose (w/o replacement) i features
 - B. Choose best of these i as the root of this (sub)tree
3. Do NOT prune

1. about choosing the features, increasing i ,

- increases correlation among individual trees (BAD)
- accuracy of individual trees (GOOD)

therefore,

- use a tuning set to pick a good value of i

(note: tuning set is a set for picking the best hyperparameters after training)

2. overall advantages:

very fast, works for a large number of features, reduces overfitting

3. extreme RF:

main differences:

- choose random split points, rather than based on IG etc.
- test time: unweighted vote

AdaBoost

- Boosting:

deal with weak learners,

focuses on the current model **correcting** incorrect predictions

- loss

error: measure the deviation between the true and predicted values

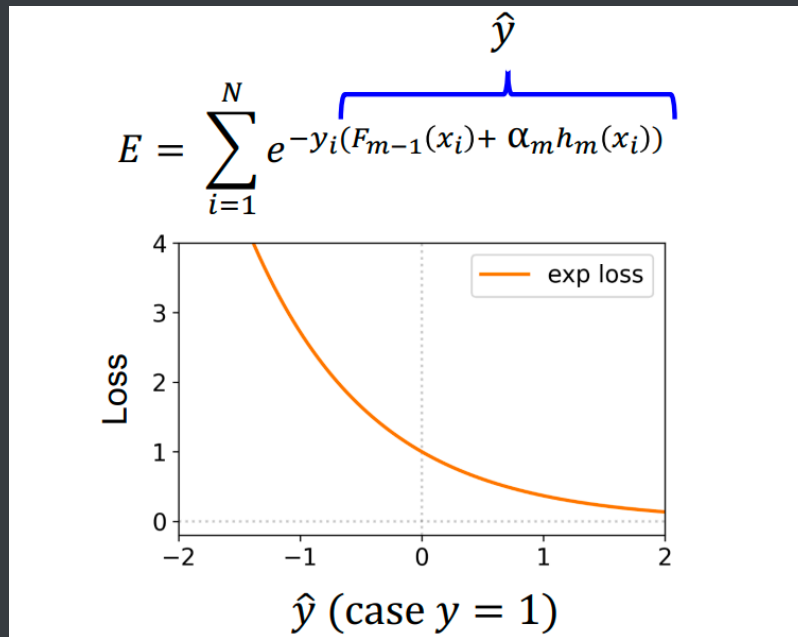
loss: quantify the impact or consequence of an error

both always > 0 , loss drops quickly

(note: the error and loss can be identical, e.g. linear regression)

in AdaBoost,

loss is small if predicted and true label have same sign (binary classification, only the sign makes a difference)



■ goal

□ **Given:** $S = \{(x_j, y_j)\}$ with $j \in \{1, \dots, n\}$ and $y \in \{-1, +1\}$

□ **Learn:** $F(X) = \alpha_1 h_1(X) + \alpha_2 h_2(X) + \dots + \alpha_t h_t(X)$

□ **Prediction:**
$$F(x_i) = \begin{cases} -1 & \text{(Negative)} \\ +1 & \text{(Positive)} \end{cases}$$

$$\text{sign}(\sum_t \alpha_t h_t(x_i))$$

□ **Goal:**

▣ Pick $\alpha_m h_m(X)$ to add to the model

▣ To minimize loss:
$$E = \sum_{i=1}^N e^{-y_i (F_{m-1}(x_i) + \alpha_m h_m(x_i))}$$

■ how to pick classifier: h_m

□ Let $w_i^{(m)} = e^{-y_i(F_{m-1}(x_i))}$ and rewrite the loss

$$= \sum_{y_i = h_m(x_i)} w_i^{(m)} e^{-\alpha_m} + \sum_{y_i \neq h_m(x_i)} w_i^{(m)} e^{\alpha_m}$$

Weight for correct predictions

Weight for incorrect predictions

$$= \sum_{i=1}^N w_i^{(m)} e^{-\alpha_m} + \sum_{y_i \neq h_m(x_i)} w_i^{(m)} (e^{\alpha_m} - e^{-\alpha_m})$$

Assumes h_m 's predictions are all correct

h_m that minimizes this sum, minimizes E!

jesse davis

■ how to pick weight: α_m

□ Let $W_c = \sum_{y_i = h_m(x_i)} w_i^{(m)}$ and $W_{ic} = \sum_{y_i \neq h_m(x_i)} w_i^{(m)}$

□ Then $E = W_c e^{-\alpha_m} + W_{ic} e^{\alpha_m}$

$$\frac{dE}{d\alpha} = -W_c e^{-\alpha_m} + W_{ic} e^{\alpha_m} = 0$$

$$-W_c + W_{ic} e^{2\alpha_m} = 0$$

$$\alpha_m = \frac{1}{2} \ln \frac{W_c}{W_{ic}}$$

$$\alpha_m = \frac{1}{2} \ln \frac{1 - \varepsilon_m}{\varepsilon_m}$$

$$\text{with } \varepsilon_m = \frac{\sum_{y_i \neq h_m(x_i)} w_i^{(m)}}{\sum w_i^{(m)}}$$

jesse davis

■ complete algorithm

Given $S = \{(x_j, y_j)\}$ where $j \in \{1, \dots, n\}$, Integer T

$$w_i^{(1)} = 1/n$$

All examples
have same weight

for $t = 1$ **to** T :

Weighted error

Find classifier h_t , with small error $\epsilon_t = \sum_{h_t(x_i) \neq y_i} w_i^{(t)}$

if $(\epsilon_t > 1/2)$ **then** break

$$\beta_t = \epsilon_t / (1 - \epsilon_t)$$

Down weight correct predictions

for $i = 1$ **to** n : **if** $(h_t(x_i) = y_i)$ **then** $w_i^{(t+1)} = w_i^{(t)} \beta_t$

for $i = 1$ **to** n : $w_i^{(t+1)} = \frac{w_i^{(t+1)}}{\sum w_j^{(t+1)}}$

Normalize weights

$$\alpha_t = \ln \frac{1}{\beta_t}$$

vis

There will be T classifiers.

- important practical details

- typically use depth bounded decision tree
- for weighted instances:

solution 1:

sample a larger set of unweighted instances,

according to the **weight distribution**,

and run learner on this new dataset

solution 2:

adapt learner,

e.g. weighted counts for split criteria for decision trees

基尼指数的加权版本

基尼指数公式:

$$G = \sum_i p(i)(1 - p(i))$$

在引入样本权重后, 每个样本的贡献变为:

$$G_{weighted} = \sum_i w_i \cdot p(i) \cdot (1 - p(i))$$

- training rounds & margins

$\text{margin} = (\text{正类的加权票数}) - (\text{负类的加权票数})$, 表示信心 (置信度) ;

AdaBoost 不容易过拟合, 更多轮的训练会增加 margin。大 margin 意味着弱分类器之间更一致, 预测结果更有信心。模型泛化能力强;

4. AdaBoost (only binary classification) variations

variations for multiclass tasks & regression

Manipulating inputs / outputs

input

different learners see different subsets of features

but does not work that well

output

- multiclass problem

solution 1:

one-vs-rest strategy: learn k binary classifiers

solution 2:

$\log k$ models

- error-correcting codes:**

construct new labels for each class

- Build T bit code word for each class
- For class y_k , i^{th} bit
 - 1 if y_k is in new 'pos' class for h_i
 - 0 if y_k is in new 'neg' class for h_i
- For $t = 1$ to T
 - Partition labels into two disjoint label sets
 - Train model on new label set

Y	Code Words				
	1	2	3	4	5
0	0	0	1	1	0
1	0	1	1	0	1
2	0	0	0	0	1
3	0	1	0	1	0
4	1	0	1	0	1
5	1	1	1	0	0
6	1	0	0	1	0
7	1	1	0	1	1

for a test result (a T-bit vector), use hamming distance to find the closest class

why ensembles work

bias / variance explanation

the expected error:

- Assume that the true function is $f(x)$

$$\begin{aligned} E_p[(y - h(x))^2] &= (h(x) - f(x))^2 && (\text{bias})^2 \\ &+ E_p[(h(x) - \overline{h(x)})^2] && \text{variance} \\ &+ E_p[(y - f(x))^2] && \text{noise} \end{aligned}$$

- ▣ **Bias:** Inability to represent the true target concept
- ▣ **Variance:** Fluctuations due to variations in data sample
- ▣ **Inherent error:** Inability to distinguish between two objects with different labels

- bias: underfitting the data

incorrect assumptions,

inability to represent certain decision boundaries (wrong model),

"too global" classifiers (e.g. small decision tree)

- variance: overfitting the data

"too local" classifiers (knn),

decision made on small subsets of data,

randomization in the learning algorithm,

unstable learning algorithm

- for ensembles

- bagging (sample)

theoretically, reduces variance

actually, resolves both

- boosting (weight)

early iterations -> bias

later iterations -> variance

statistical explanation

- 贝叶斯后验估计
- bayes optimal classifier:

AdaBoost、Random Forest 等方法通过将多个弱分类器的预测结果进行加权投票，从而模拟贝叶斯分类器的后验概率推断过程。

- ensembles 近似 bayes optimal.

representational explanation

线性加权平均能够更好地逼近真实函数

computational explanation

avoid local minima -> repeat the search many times with random restarts

Applications