


# experimental methodology

## generalizing

- test set, training set
- cross validation: make the most of data



example:

| Training Folds | Testing Fold | Correct | Incorrect | Accuracy |
|----------------|--------------|---------|-----------|----------|
| S1, S2, S3, S4 | S5           | 35      | 15        | 0.70     |
| S1, S2, S3, S5 | S4           | 30      | 20        | 0.60     |
| S1, S2, S4, S5 | S3           | 37      | 13        | 0.74     |
| S1, S3, S4, S5 | S2           | 33      | 17        | 0.66     |
| S2, S3, S4, S5 | S1           | 40      | 10        | 0.80     |

but one class might be completely missing in the training data

problem: an infrequent class might be missing at all

-> solution: stratification, then concatenate all test results

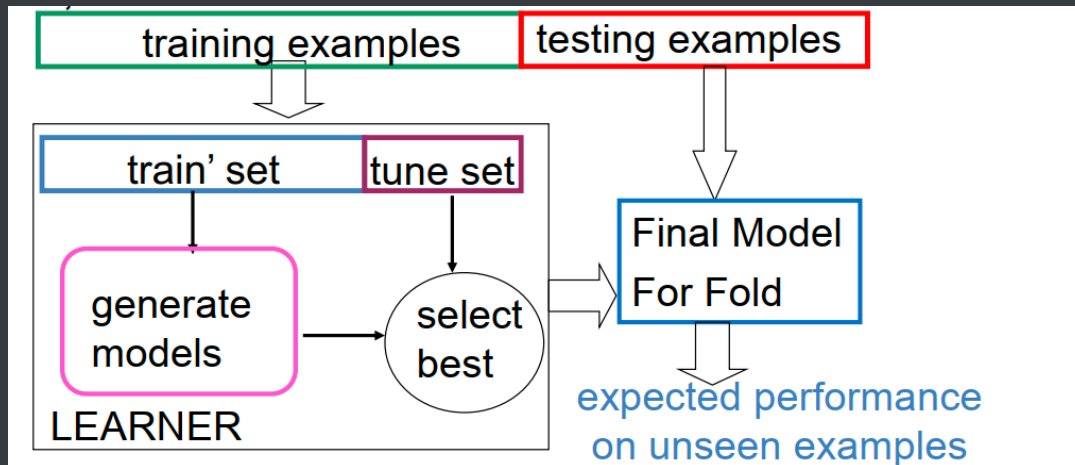
- leave-one-out cross validation

leave only one but not a set out

very computationally expensive

- big pitfall: test data is not necessarily representative of real data!

# Tuning



- internal cross-validation

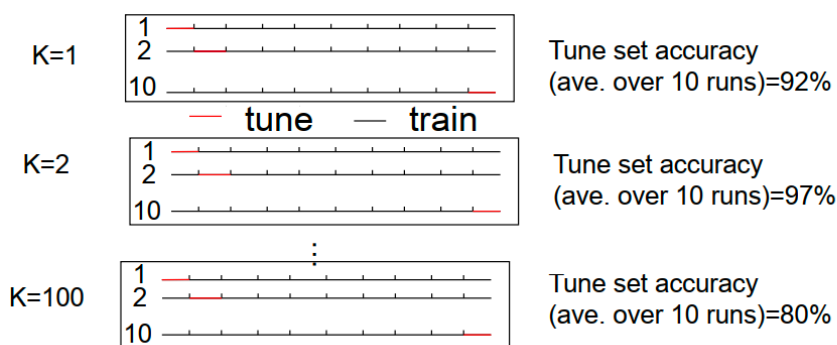
motivation: single tuning set can be unreliable

- CV on training data to select the parameters
- Score setting by average performance
- Pick best one
- Train model on all training data

example:

Step 1: Try various values for  $k$  in  $k$ -NN

Use 10 train/tune splits for each  $k$



Step 2: Select best value for  $k$

Step 3: Learn model using all training data

Step 4: Make predictions on test data

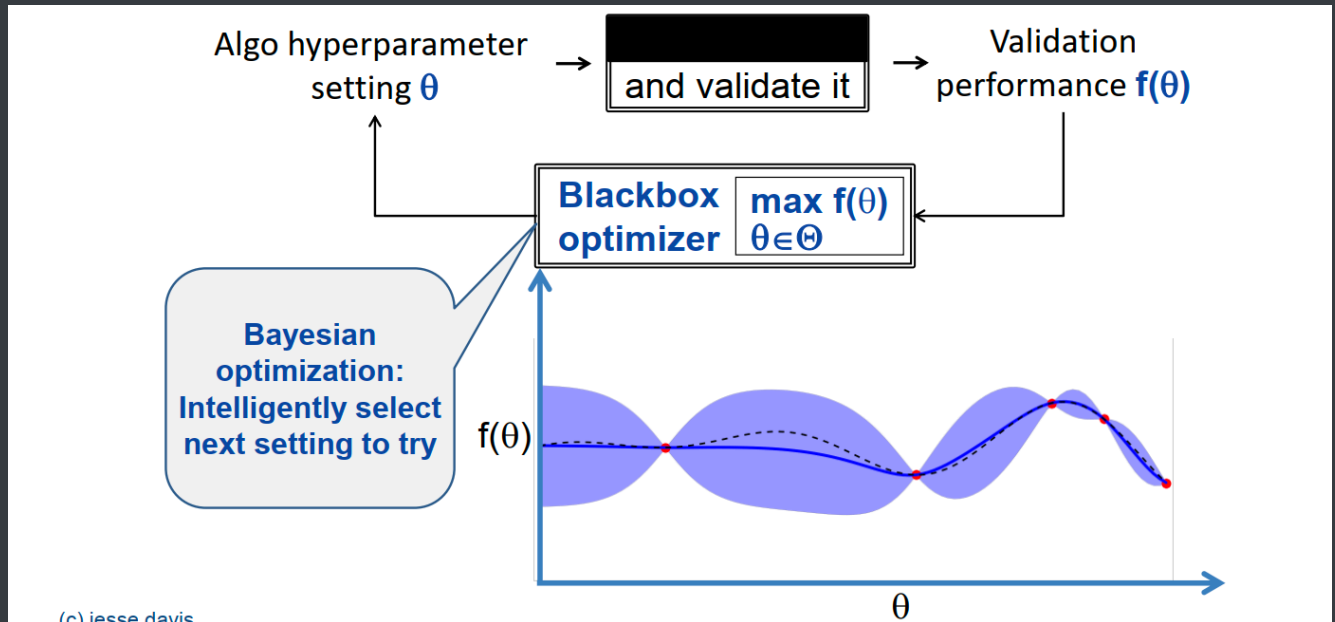
- spread of values: wide-range, 避免等差数列 (tune之前的参数设置需要拉开差距)
- hyperparameter search

search problem

random search (sample possible values) <-> grid search

random search 通常很有效，因为存在一些效果差不多的configurations，grid search效率很低

### automated hyperparameter tuning



红色的点是采样点，第二个紫色包中可能存在更好的点

#### parameter-less algorithm:

internally sets all parameters, rather than user explicitly setting the hyperparameters

to wrap up,

## Deploying a System

26

- Use 10 fold CV to pick best algorithm
  - ▣ Tune each algorithm (parameters, feature sets, etc.)
  - ▣ Give performance estimate of best approach to user
- For best algorithm, only use tuning sets to pick settings (features, parameters)
- Train model using all the data and selected settings and deploy this model

## comparisons

- 假设检验 hypothesis testing:

null hypothesis 原假设  $H_0$  - 两个模型performance没有显著差异

alternative hypothesis 备选假设  $H_a$  - 两个模型performance有显著差异


计算统计量和p值, 比较p值和显著性水平;

### Two Tailed test

$H_0$ : No difference in model performance

$H_a$ : Model's perform differently

P-value = **the probability** of your data or something more extreme under the null hypothesis.



- McNemar's Test

core: how often was  $f_1$  right and  $f_2$  wrong on the sample example, or the other way around?

|             | $f_1$<br>right | $f_1$<br>wrong |
|-------------|----------------|----------------|
| $f_2$ right | A              | B              |
| $f_2$ wrong | C              | D              |

to see if B and C are significantly different:

$$\chi^2 = \frac{(b - c)^2}{(b + c)},$$

- paired t-test

|            | <u>Accuracies on Testsets</u> |    |    |     |    |
|------------|-------------------------------|----|----|-----|----|
| Algo L1:   | 88                            | 69 | 79 | ... | 72 |
| Algo L2:   | 79                            | 62 | 74 | ... | 75 |
| $\delta$ : | +9                            | +7 | +5 | ... | -3 |


null hypothesis: algos have equivalent average accuracies

if p-value is low, then reject null hypothesis.

assumptions of the t-test:

1. test statistic is normally distributed

e.g. if metric is classifier accuracy, then 

AUC ROC then  because usually not normally distributed

-> wilcoxon signed-rank test: 非参数检验方法，不需要正态性假设

2. independent sample of test-examples

cross validation通常违反这一条规则，但实际上t检验对独立性假设的偏差是非常robust的；

#### ■ sign test

记录两个算法L1和L2在测试样本上“意见不一致”情况下的“win”次数，M是较大的win次数；

假设两个模型性能相似，用二项分布 $b(N, 0.5)$ 计算L1或者L2在随机情况下赢至少M次的概率，如果概率很低则拒绝假设；

#### ■ wilcoxon signed rank test - on multiple data sets

|  | RF | DT | RF-DT | RF-DT | Rank | Signed Rank |
|--|----|----|-------|-------|------|-------------|
|  | 90 | 88 | 2     | 2     | 2    | 2           |
|  | 99 | 66 | 33    | 33    | 10   | 10          |
|  | 65 | 75 | -10   | 10    | 6    | -6          |
|  | 67 | 55 | 12    | 12    | 7    | 7           |
|  | 87 | 73 | 14    | 14    | 8    | 8           |
|  | 69 | 75 | -6    | 6     | 4    | -4          |
|  | 73 | 72 | 1     | 1     | 1    | 1           |
|  | 74 | 89 | -15   | 15    | 9    | -9          |
|  | 80 | 85 | -5    | 5     | 3    | -3          |
|  | 77 | 70 | 7     | 7     | 5    | 5           |

$W_+ = 33$  and  $W_- = 22$   
 $T_{\text{Wilcox}} = \min(33, 22) = 22$

calculation:

step 1 compute performance difference

step 2 rank the absolute differences


step 3 add sign to rank

step 4 sum of positive ranks & sum of negative ranks, take the minimum value

meaning:

generally W值越大说明两个模型差异越大，但需要和P值配合来跟显著性水平比较！

■ graphical comparison



一个数据点代表一个数据集上的结果，黑色直线代表两个模型性能相同