

reinforcement learning : intro + MDP

terminologies

actions, probabilities, costs (of actions), utilities (of results)

decision-theoretic agent

an agent is rational exactly when it chooses the action with the maximum expected utility taken over all results of actions

utility-theoretic agent

Expected Utility:

$$EU(A \mid E) = \sum_i P(Result_i(A) \mid Do(A), E) U(Result_i(A))$$

The **principle of maximum expected utility (MEU)** says that a rational agent should choose an action that maximizes $EU(A \mid E)$.

axioms of utility

a lottery:

Example:

Lottery L with two outcomes, C_1 and C_2 :

$$L = [p, C_1; 1 - p, C_2]$$

Preference between lotteries:

$L_1 \succ L_2$ The agent prefers L_1 over L_2

$L_1 \sim L_2$ The agent is indifferent between L_1 and L_2

$L_1 \succsim L_2$ The agent prefers L_1 or is indifferent between L_1 and L_2

the axioms of utility

- **Orderability**

$$(A \succ B) \vee (B \succ A) \vee (A \sim B)$$

- **Transitivity**

$$(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$$

- **Continuity**

$$A \succ B \succ C \Rightarrow \exists p[p, A; 1 - p, C] \sim B$$

- **Substitutability**

$$A \sim B \Rightarrow [p, A; 1 - p, C] \sim [p, B; 1 - p, C]$$

- **Monotonicity**

$$A \succ B \Rightarrow (p > q \Leftrightarrow [p, A; 1 - p, B] \succ [q, A; 1 - q, B])$$

- **Decomposability**

$$[p, A; 1 - p, [q, B; 1 - q, C]] \sim [p, A; (1 - p)q, B; (1 - p)(1 - q), C]$$

utility functions and axioms

- **Utility Principle** If an agent's preferences obey the axioms, then there exists a function $U : S \mapsto R$ with
 $U(A) > U(B) \Leftrightarrow A \succ B$
 $U(A) = U(B) \Leftrightarrow A \sim B$

- **Expected Utility of a Lottery:**

$$U([p_1, S_1; \dots; p_n, S_n]) = \sum_i p_i U(S_i)$$

→ Since the outcome of a nondeterministic action is a lottery, an agent can act rationally only by following the Maximum Expected Utility (MEU) principle.

assign utilities

- "Best possible prize" $U(S) = u_{max} = 1$
- "Worst catastrophe" $U(S) = u_{min} = 0$

Given a utility scale between u_{min} and u_{max} we can assess the utility of any particular outcome S by asking the agent to choose between S and a standard lottery $[p, u_{max}; 1 - p, u_{min}]$. We adjust p until they are equally preferred.

Then, p is the utility of S . This is done for each outcome S to determine $U(S)$.

human judgement and irrationality

sequential decision problems

deterministic ~: all actions always lead to the next selected direction

stochastic ~: each action achieves the intended effect with a probability

- an MDP:

- **Set of states** S
- **Set of actions** A
- **Transition model** $P(s' | s, a)$, with $s, s' \in S$ and $a \in A$
- **Reward function** $R(s)$, with $s \in S$ **ALTERNATIVE is to use $R(s, a, s')$**

- transition model, policy, optimal policy
- performance: sum of rewards for the states visited
- finite / infinite horizon

finite horizon: 有限时间范围, 到达一个 s_N 之后的utility就不再重要

- Finite horizon: $U_h([s_0, s_1, \dots, s_{N+k}]) = U_h([s_0, s_1, \dots, s_N])$ for all $k > 0$.

optimal policy depends on **current state & remaining steps to go**

(nonstationary)

for infinite horizon:

optimal policy only depends on current state (stationary)

- assign utilities:

- **Additive rewards:**
 $U_h([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$
- **Discounted rewards:**
 $U_h([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$
- The term $\gamma \in [0, 1[$ is called the **discount factor**.
- With **discounted rewards** the **utility of an infinite state sequence is always finite**. The **discount factor expresses** that **future rewards have less value than current rewards**.

for finite horizon, additive rewards are reasonable;

for infinite horizon, discount \rightarrow reward value converges to a finite value

- the utility of a state depends on the utility of the state sequences that follow it

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t) \mid \pi, s_0 = s \right]$$

Note:

$R(s)$ is the short-term reward for being in s

$U(s)$ is the long-term total reward from s onwards

- choosing actions using MEU

$$\pi(s) = \operatorname{argmax}_a \sum_{s'} P(s' \mid s, a) U(s')$$

Bellman-equation:

$$U(s) = R(s) + \gamma \max_a \sum_{s'} P(s' \mid s, a) U(s')$$

utility of a state =

immediate reward

+ expected discounted utility of the next state

应用bellman equation的例子:

Bellman-Equation: Example

- In our 4×3 world the equation for the state (1,1) is

$$\begin{aligned}
 U(1, 1) &= -0.04 + \gamma \max \{ \begin{array}{l} 0.8U(1, 2) + 0.1U(2, 1) + 0.1U(1, 1), \quad (Up) \\ 0.9U(1, 1) + 0.1U(1, 2), \quad (Left) \\ 0.9U(1, 1) + 0.1U(2, 1), \quad (Down) \\ 0.8U(2, 1) + 0.1U(1, 2) + 0.1U(1, 1) \} \quad (Right) \\
 &= -0.04 + \gamma \max \{ \begin{array}{l} 0.8 \cdot 0.762 + 0.1 \cdot 0.655 + 0.1 \cdot 0.705, \quad (Up) \\ 0.9 \cdot 0.705 + 0.1 \cdot 0.762, \quad (Left) \\ 0.9 \cdot 0.705 + 0.1 \cdot 0.655, \quad (Down) \\ 0.8 \cdot 0.655 + 0.1 \cdot 0.762 + 0.1 \cdot 0.705 \} \quad (Right) \\
 &= -0.04 + 1.0 (0.6096 + 0.0655 + 0.0705), \quad (Up) \\
 &= -0.04 + 0.7456 = 0.7056
 \end{aligned}$$

for these utilities

- Up is the optimal action in (1,1).

3	0.812	0.868	0.918	<div>+1</div>
2	0.762		0.660	<div>-1</div>
1	0.705	0.655	0.611	0.388
	1	2	3	4

注意stochastic agent的特点：预期方向和预期方向的正交方向！

- q-equation: generalization of the Bellman-equation

$$Q(s, a) = R(s) + \gamma \sum_{s'} P(s' | s, a) U(s')$$

$$Q(s, a) = R(s) + \gamma \sum_{s'} P(s' | s, a) \max_{a'} Q(s', a')$$

is the corresponding equation for the **Q-function -the action utility function** - the utility after taking action a in s

$$U(s) = \max_a Q(s, a)$$

value iteration: calculating optimal policies

histories, separable,

A utility function on histories U_h is **separable** iff there exists a function f such that

$$U_h([s_0, s_1, \dots, s_n]) = f(s_0, U_h([s_1, \dots, s_n]))$$

The simplest form is an additive reward function R :

$$U_h([s_0, s_1, \dots, s_n]) = R(s_0) + U_h([s_1, \dots, s_n])$$

-> 递归迭代（如果计算过程有环，history的长度可能是无限长；而无限长也可能最终值是收敛的converge）

非线性（max运算）

算法：（多次迭代直到收敛）

```
function VALUE-ITERATION( $mdp, \epsilon$ ) returns a utility function
  inputs:  $mdp$ , an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ ,
           rewards  $R(s, a, s')$ , discount  $\gamma$ 
            $\epsilon$ , the maximum error allowed in the utility of any state
  local variables:  $U, U'$ , vectors of utilities for states in  $S$ , initially zero
                      $\delta$ , the maximum relative change in the utility of any state

  repeat
     $U \leftarrow U'; \delta \leftarrow 0$ 
    for each state  $s$  in  $S$  do
       $U'[s] \leftarrow \max_{a \in A(s)} Q\text{-VALUE}(mdp, s, a, U)$ 
      if  $|U'[s] - U[s]| > \delta$  then  $\delta \leftarrow |U'[s] - U[s]|$ 
  until  $\delta \leq \epsilon(1 - \gamma)/\gamma$ 
  return  $U$ 
```

为什么这样能收敛？可证明：

- It can be shown that **value iteration converges** and that

$$\text{if } \|U_{t+1} - U_t\| < \epsilon(1 - \gamma)/\gamma \quad \text{then} \quad \|U_{t+1} - U\| < \epsilon$$

$$\text{if } \|U_t - U\| < \epsilon \quad \text{then} \quad \|U^{\pi_t} - U\| < 2\epsilon\gamma/(1 - \gamma)$$

policy iteration

- policy evaluation

- **Policy evaluation:** given a policy π_t , calculate $U_t = U^{\pi_t}$, the utility of each state if π_t were executed.

$$U_i(s) = R(s) + \gamma \sum_{s'} P(s'|s, \pi_i(s)) U_i(s')$$

- policy improvement

$$\pi_{t+1}(s) = \operatorname{argmax}_a \sum_{s'} P(s' | s, a) U_t(s')$$

算法伪代码：

```

function POLICY-ITERATION(mdp) returns a policy
  inputs: mdp, an MDP with states S, actions A(s), transition model  $P(s' | s, a)$ 
  local variables: U, a vector of utilities for states in S, initially zero
                    $\pi$ , a policy vector indexed by state, initially random

  repeat
    U  $\leftarrow$  POLICY-EVALUATION( $\pi$ , U, mdp)
    unchanged?  $\leftarrow$  true
    for each state s in S do
       $a^* \leftarrow \operatorname{argmax}_{a \in A(s)} \text{Q-VALUE}(\text{mdp}, s, a, U)$ 
      if Q-VALUE(mdp, s,  $a^*$ , U) > Q-VALUE(mdp, s,  $\pi[s]$ , U) then
         $\pi[s] \leftarrow a^*$ ; unchanged?  $\leftarrow$  false
  until unchanged?
  return  $\pi$ 

```

多次迭代对每一个state都找到最佳action（每个state的最佳action都left unchanged）

value iteration 和 policy iteration 两个算法的区别

policy iteration算法实际上是value iteration算法的一部分，value iteration直接计算value