

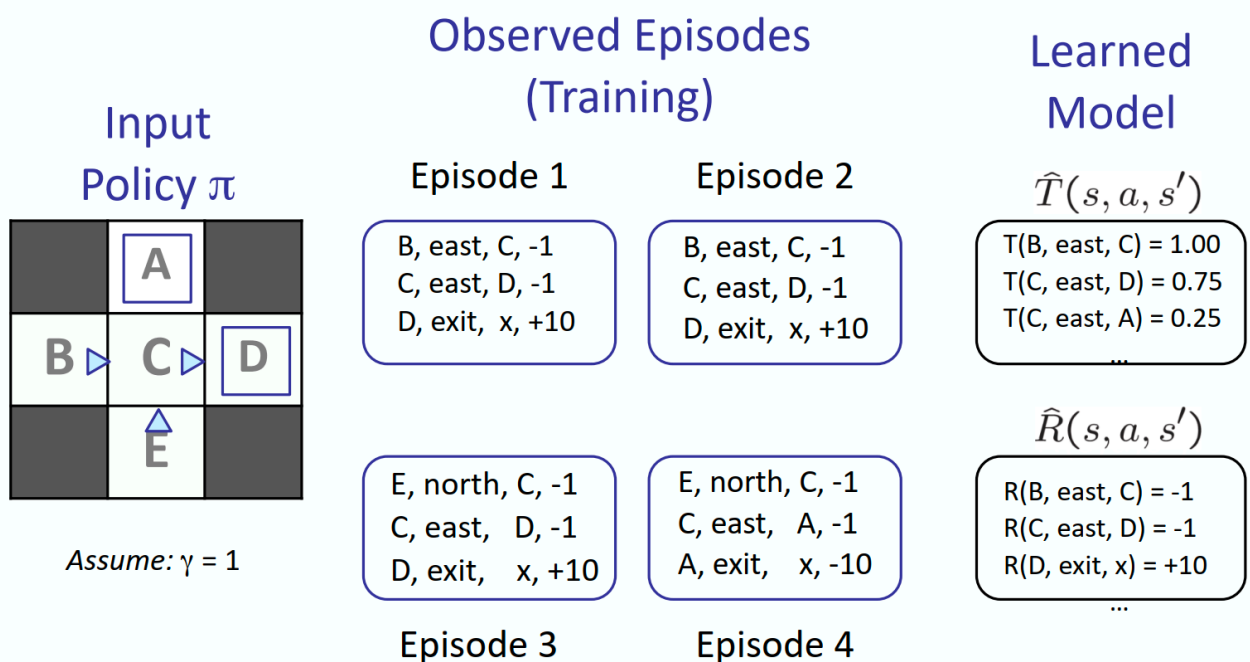
Reinforcement learning (1)

model-based

- Still assume a Markov decision process (MDP):
 - A **set of states** $s \in S$
 - A **set of actions** (per state) A
 - A **model** $T(s,a,s')$
 - A **reward function** $R(s,a,s')$
- Still looking for a policy $\pi(s)$
- New twist: **don't know T or R**

- idea: learn empirical MDP model \rightarrow solve the learned MDP
- example

Example: Model-Based Learning



- model-based是先学一个完整model (T和R) , model-free是绕过model

model-free

passive reinforcement learning (value learning)

- task: policy evaluation

given a fixed policy (still don't know T/R), learn the state values

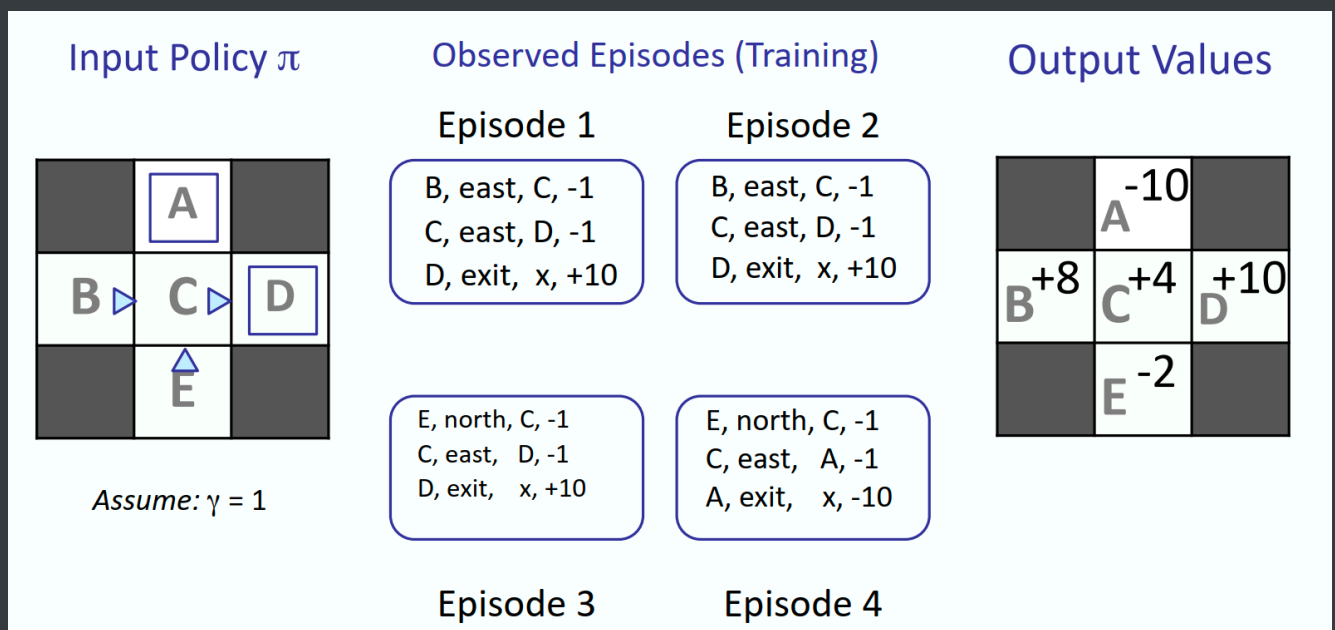
- learner just executes the policy and learns from experience!

this is NOT offline planning!

offline planning 是已知T和R, 用value iteration或policy iteration的方式提前规划出最优行动

- direct evaluation

average observed sample values



easy to perform, but wastes information about state connections

- temporal difference learning

learn from every experience:

Sample of $V(s)$: $sample = R(s, \pi(s), s') + \gamma V^\pi(s')$

Update to $V(s)$: $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + (\alpha)sample$

makes recent samples more important; alpha is learning rate

decreasing learning rate can give converging averages

problem: not able to turn values into a better policy

the idea would be, to learn Q-values rather than values (you choose the actions!)

active reinforcement learning (q-learning)

- full reinforcement learning: don't know transitions / rewards, choose the actions to learn the optimal policy / values
- recap: q-value iteration

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma \max_{a'} Q_k(s', a') \right]$$


- q-learning

- Learn $Q(s,a)$ values as you go
 - Receive a sample (s,a,s',r)
 - Consider your old estimate: $Q(s,a)$
 - Consider your new sample estimate:

$$sample = R(s, a, s') + \gamma \max_{a'} Q(s', a')$$

- Incorporate the new estimate into a running average:

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha) [sample]$$



q-learning converges to optimal policy, even the actions are suboptimal;

but, explore enough, eventually the learning rate should be small enough

summary of the algorithms

Known MDP: Offline Solution

Goal

Compute V^*, Q^*, π^*

Evaluate a fixed policy π

Technique

Value / policy iteration

Policy evaluation

Unknown MDP: Model-Based

Goal

Compute V^*, Q^*, π^*

Evaluate a fixed policy π

Technique

VI/PI on approx. MDP

PE on approx. MDP

Unknown MDP: Model-Free

Goal

Compute V^*, Q^*, π^*

Evaluate a fixed policy π

Technique

Q-learning

Value Learning

exploration vs. exploitation

- random actions

- Simplest: random actions (ϵ -greedy)
 - Every time step, flip a coin
 - With (small) probability ϵ , act randomly
 - With (large) probability $1-\epsilon$, act on current policy

problem:

keep thrashing around once learning is done

(掌握了最优policy之后依然有epsilon的概率随机选择动作)

solution 1: lower epsilon over time

solution 2: exploration functions

- exploration functions

basic idea:

introduce a visit count, explore not-enough-visited areas with higher priority

- Takes a value estimate u and a visit count n , and returns an optimistic utility, e.g. $f(u, n) = u + k/n$

when n is small, f would be larger

when n is large, f would be approaching the q value

prioritized q-learning update function:

Modified Q-Update: $Q(s, a) \leftarrow_{\alpha} R(s, a, s') + \gamma \max_{a'} f(Q(s', a'), N(s', a'))$

bonus propagation:

if s' is an unknown state, then $Q(s', a')$ would be large, which would be propagated to $Q(s, a)$

- regret

衡量学习过程中的错误代价 (we wanna optimally learn to be optimal)

random exploration has higher regret than exploration functions

approximate q-learning

in realistic situations we cannot possibly learn about every single state

-> generalization is important

- feature-based representations

a vector,

e.g. in pacman features might be distance to closest ghost / dot etc.

- linear value functions

- Using a feature representation, we can write a q function (or value function) for any state using a few weights:

$$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- approximate Q-learning

- Q-learning with linear Q-functions:

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

transition = (s, a, r, s')

$$\text{difference} = \left[r + \gamma \max_{a'} Q(s', a') \right] - Q(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$$

Exact Q's

$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$$

Approximate Q's

- reasoning: minimizing error

$$\text{error}(w) = \frac{1}{2} \left(y - \sum_k w_k f_k(x) \right)^2$$

$$\frac{\partial \text{error}(w)}{\partial w_m} = - \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$

$$w_m \leftarrow w_m + \alpha \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$

Approximate q update explained:

$$w_m \leftarrow w_m + \alpha \left[\underbrace{r + \gamma \max_{a'} Q(s', a')}_{\text{"target"}} - \underbrace{Q(s, a)}_{\text{"prediction"}} \right] f_m(s, a)$$

"target"

"prediction"

