

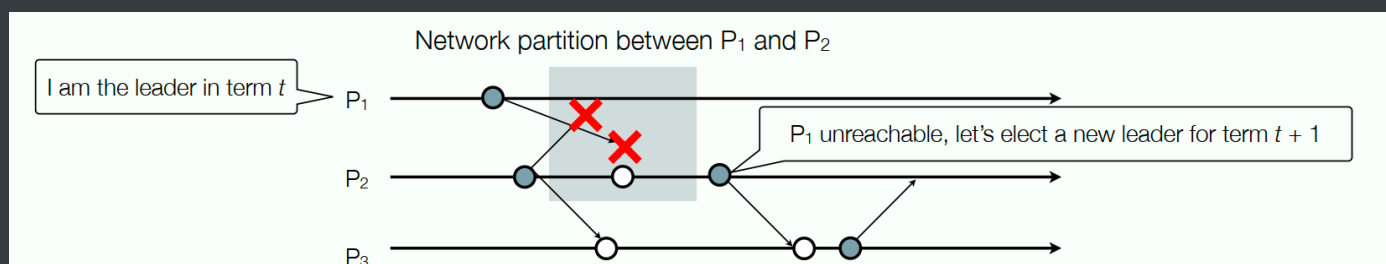
# consensus – 2

## leader election

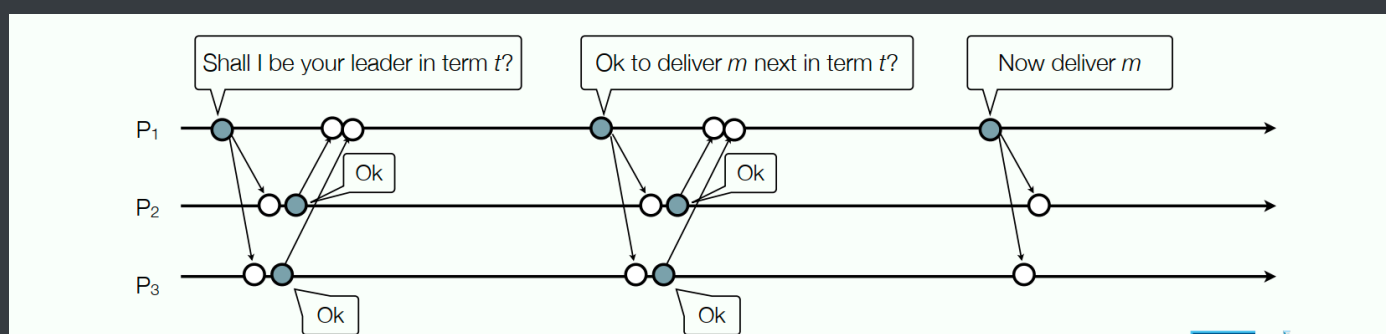
- time-out failure detector,  
on suspected leader crash, elect a new one
- prevent 2 leaders at the same time
- **term**: incremented once a leader election is started  
one node can only vote once per term  
require a quorum of nodes to elect a leader in a term

## one-leader guarantee

如果leader失联，则elect一个新leader：



old leader无法越权，因为消息发送前需要获得一个quora统一，而old leader至慢也会在第二轮ack中得知自己不再是leader



## consensus algorithms

- viewstamped replication
- paxos

agreement on a single value (multi-paxos is for a sequence of value)

Google Chubby

- raft

specifically for log replication

supports sequences of values

## the Raft consensus algorithm

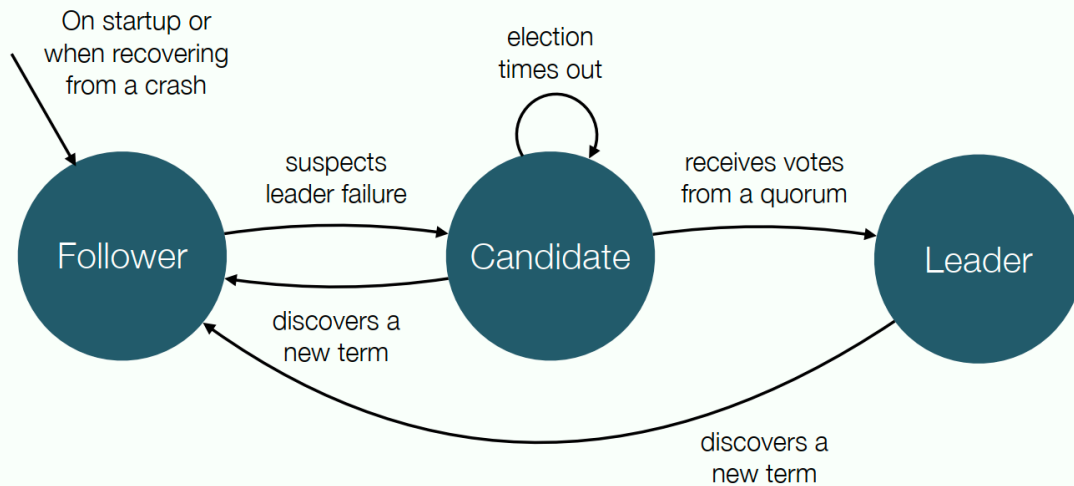
### system model

- partially synchronous, crash-recovery
  - uses clocks only for **liveness**
- processes may fail, network may fail
- **cooperative processes**, does not deal with malicious ("byzantine") processes

### basic idea

- term is an integer counter
- **consistent logs**
  - one leader per term, all other nodes are followers
  - ordered entries, messages + terms
  - only the leader can append new entries to the log**
  - (log is append-only !!! grow only at the end)*

### algorithm



- initialisation

```

on initialisation do
  currentTerm := 0; votedFor := null
  log := {}; commitLength := 0
  currentRole := follower; currentLeader := null
  votesReceived := {}; sentLength := {}; ackedLength := {}
end on

on recovery from crash do
  currentRole := follower; currentLeader := null
  votesReceived := {}; sentLength := {}; ackedLength := {}
end on

```

- currentTerm, votedFor, log, commitLength are stored in persistent memory
- currentRole, currentLeader, votesReceived, sentLength, ackedLength can be stored in volatile memory, which could be reset during crash-recovery

- starting a leader election

```

on node nodeId suspects leader has failed, or on election timeout do
  currentTerm := currentTerm + 1; currentRole := candidate
  votedFor := nodeId; votesReceived := {nodeId}; lastTerm := 0
  if log.length > 0 then lastTerm := log[log.length - 1].term; end if
  msg := (VoteRequest, nodeId, currentTerm, log.length, lastTerm)
  for each node in nodes: send msg to node
  start election timer
end on

```

votes for itself,

sends a VoteRequest msg to each other node

sets a timer, if times out then steps are repeated

- voting on a new leader

```

on receiving (VoteRequest, cId, cTerm, cLogLength, cLogTerm)
  at node nodeId do
  if cTerm > currentTerm then
    currentTerm := cTerm; currentRole := follower
    votedFor := null
  end if
  lastTerm := 0
  if log.length > 0 then lastTerm := log[log.length - 1].term; end if
  logOk := (cLogTerm > lastTerm) ∨
    (cLogTerm = lastTerm ∧ cLogLength ≥ log.length)

  if cTerm = currentTerm ∧ logOk ∧ votedFor ∈ {cId, null} then
    votedFor := cId
    send (VoteResponse, nodeId, currentTerm, true) to node cId
  else
    send (VoteResponse, nodeId, currentTerm, false) to node cId
  end if
end on

```

## ■ collecting votes

```

on receiving (VoteResponse, voterId, term, granted) at nodeId do
  if currentRole = candidate ∧ term = currentTerm ∧ granted then
    votesReceived := votesReceived ∪ {voterId}
    if |votesReceived| ≥ ⌈(|nodes| + 1)/2⌉ then
      currentRole := leader; currentLeader := nodeId
      cancel election timer
      for each follower ∈ nodes \ {nodeId} do
        sentLength[follower] := log.length
        ackedLength[follower] := 0
        REPLICATELOG(nodeId, follower)
      end for
    end if
  else if term > currentTerm then
    currentTerm := term
    currentRole := follower
    votedFor := null
    cancel election timer
  end if
end on

```

## ■ broadcasting messages

```

on request to broadcast msg at node nodeId do
  if currentRole = leader then
    append the record (msg : msg, term : currentTerm) to log
    ackedLength[nodeId] := log.length
    for each follower ∈ nodes \ {nodeId} do
      REPLICATELOG(nodeId, follower)
    end for
  else
    forward the request to currentLeader via a FIFO link
  end if
end on

periodically at node nodeId do
  if currentRole = leader then
    for each follower ∈ nodes \ {nodeId} do
      REPLICATELOG(nodeId, follower)
    end for
  end if
end do

```

## ■ replicating from leader to followers

```

function REPLICATELOG(leaderId, followerId)
    prefixLen := sentLength[followerId]
    suffix := ⟨log[prefixLen], log[prefixLen + 1], ...,
              log[log.length - 1]⟩
    prefixTerm := 0
    if prefixLen > 0 then
        prefixTerm := log[prefixLen - 1].term
    end if
    send (LogRequest, leaderId, currentTerm, prefixLen,
          prefixTerm, commitLength, suffix) to followerId
end function

```

- followers receiving messages from a leader

```

on receiving (LogRequest, leaderId, term, prefixLen, prefixTerm,
              leaderCommit, suffix) at node nodeId do
    if term > currentTerm then
        currentTerm := term; votedFor := null
        cancel election timer
    end if
    if term = currentTerm then
        currentRole := follower; currentLeader := leaderId
    end if
    logOk := (log.length ≥ prefixLen) ∧
              (prefixLen = 0 ∨ log[prefixLen - 1].term = prefixTerm)
    if term = currentTerm ∧ logOk then
        APPENDENTRIES(prefixLen, leaderCommit, suffix)
        ack := prefixLen + suffix.length
        send (LogResponse, nodeId, currentTerm, ack, true) to leaderId
    else
        send (LogResponse, nodeId, currentTerm, 0, false) to leaderId
    end if
end on

```

- updating the followers' logs

```

function APPENDENTRIES(prefixLen, leaderCommit, suffix)
    if suffix.length > 0 ∧ log.length > prefixLen then
        index := min(log.length, prefixLen + suffix.length) - 1
        if log[index].term ≠ suffix[index - prefixLen].term then
            log := ⟨log[0], log[1], ..., log[prefixLen - 1]⟩
        end if
    end if
    if prefixLen + suffix.length > log.length then
        for i := log.length - prefixLen to suffix.length - 1 do
            append suffix[i] to log
        end for
    end if
    if leaderCommit > commitLength then
        for i := commitLength to leaderCommit - 1 do
            deliver log[i].msg to the application
        end for
        commitLength := leaderCommit
    end if
end function

```

- leader receiving log acknowledgements

```

on receiving (LogResponse, follower, term, ack, success) at nodeId do
  if term = currentTerm  $\wedge$  currentRole = leader then
    if success = true  $\wedge$  ack  $\geq$  ackedLength[follower] then
      sentLength[follower] := ack
      ackedLength[follower] := ack
      COMMITLOGENTRIES()
    else if sentLength[follower] > 0 then
      sentLength[follower] := sentLength[follower] - 1
      REPLICATELOG(nodeId, follower)
    end if
  else if term > currentTerm then
    currentTerm := term
    currentRole := follower
    votedFor := null
    cancel election timer
  end if
end on

```

- leader committing log entries


```

define acks(length) =  $|\{n \in \text{nodes} \mid \text{ackedLength}[n] \geq \text{length}\}|$ 

function COMMITLOGENTRIES
  minAcks :=  $\lceil (|\text{nodes}| + 1) / 2 \rceil$ 
  ready :=  $\{len \in \{1, \dots, \text{log.length}\} \mid \text{acks}(len) \geq \text{minAcks}\}$ 
  if ready  $\neq \{\}$   $\wedge$  max(ready) > commitLength  $\wedge$ 
    log[max(ready) - 1].term = currentTerm then
    for i := commitLength to max(ready) - 1 do
      deliver log[i].msg to the application
    end for
    commitLength := max(ready)
  end if
end function

```

## evaluating Raft

- guarantees safety (agreement & validity): consistent committed log entries  
 $\leq 1$  leader per term & msg from older terms are ignored
- cannot guarantee liveness (all log entries will eventually be committed) <- FLP impossibility, but in practice 
- timeouts for failed leader detection, randomness in election, quorum

## Byzantine fault tolerance

- CFT (crash fault-tolerant) / BFT (byzantine fault-tolerant)
- CFT: processes implement the consensus algorithm correctly,
  - CFT tolerates up to (not including) 1/2 failing processes
  - BFT tolerates up to (not including) 1/3 failing processes

- "byzantine failure":


totally arbitrary failures, including

- Failing to respond to messages
- Returning incorrect results from messages
- Returning deliberately misleading results from messages
- Returning a different result for the same request to different processes (!)

- in an **adversarial** context: attacks -> byzantine failures
- in **real-world deployments**: bugs -> byzantine failures

## the byzantine generals problem

- $N = 3f + 1$



$f$  honest but faulty processes,  $f$  traitors

then # available honest =  $N - 2f$ , need to outnumber the traitors

$$N - 2f > f \iff N > 3f$$

- $N = 3f + 1$  or  $f = \text{floor}((N-1)/3)$

$N$	Honest quorum $f+1$	Faulty or traitors $f$
1	-	0
2	-	0
3	-	0
4	2	1
5	2	1
6	2	1
7	3	2
8	3	2

最少需要4 machines才能tolerate一个traitor

实际应用数量一般是7

(对比3和5, non-byzantine的结果)

- BFT consensus algorithms exist, **complexer** then CFT alg.

<- digital signatures / cryptographic has functions,

to ensure unforgeable & irrefutable

- **blockchains**: require byzantine consensus, **open and adversarial** environment, order of transactions

## consensus vs. atomic commit protocols

atomic commit protocol: commit or abort (a transaction)?

Atomic Commit	Consensus
Every process votes whether to commit or abort	Any process may propose <i>any</i> value to agree on
Must commit if <i>all</i> processes vote to commit; must abort if <i>at least one</i> node votes to abort	Any one of the proposed values is agreed on (decided)
Must abort if <i>any</i> participating process crashes	Crashed processes can be tolerated, as long as a quorum (majority) has decided