

# transactions – 1

## introduction

- distributed systems: shared data & transactions
- transaction: sequence of operations as individual unit
- critical section  
short duration, indivisible group of instructions
- atomic operation  
isolated and free of interference from other concurrent operations
- transaction: extend critical section / atomic operation  
may contain operations on different servers  
possibly long duration

## ACID

- Atomicity  
all-or-nothing, commit or abort
- Consistency  
moves data from one consistent state to another  
transaction不会破坏数据库的完整性规则  
比如订单扣减库存，如果库存不够则transaction abort并回滚，不会出现负数库存
- Isolation  
2 parts:
  - serializability  
concurrent transactions has the same effect as some serial ordering of the transactions
  - failure isolation:  
a transaction cannot see the uncommitted effects of another transaction
- Durability

preserved once a transaction commits

## life histories

```
OpenTransaction() -> Trans
CloseTransaction(Trans) -> (Commit, Abort)
AbortTransaction(Trans)
```

- success
- abort by client `AbortTransaction(Trans)`
- abort by server (then error reported)

## concurrency

- the lost update problem

interleaved execution of operations on B?

### • The lost update problem:

Transaction T A → B: 4

`bt := A.read();`

`A.write(bt-4);`

`bt := B.read();`

`bt=200`

`B.write(bt+4);`

A: 96

B: 204

C: 297

Transaction U C → B: 3

`bu := C.read();`

`C.write(bu-3);`

`bu := B.read();`

`B.write(bu+3);`

Correct B = 207!!

- inconsistent retrievals

- The inconsistent retrieval problem:

Transaction T A → B: 50

bt := A.read();  
A.write(bt-50);

bt := B.read();  
B.write(bt+50);

A: 50

B: 250

C: 300

Correct total: 600

Transaction U BranchTotal

bu := A.read();  
bu := bu + B.read();  
bu := bu + C.read();

550

- solutions?

**concurrency control:** execute transactions in such a way that overall execution is equivalent with some serial execution

-> **serially equivalent interleaving**

- The lost update problem: **serially equivalent interleaving**

Transaction T A → B: 4

bt := A.read();  
A.write(bt-4);

bt := B.read();  
B.write(bt+4);

A: 96

B: 207

C: 297

Transaction U C → B: 3

bu := C.read();  
C.write(bu-3);

bu := B.read();  
B.write(bu+3);

## recovery

- a dirty read problem

# Transactions: Recovery

- A dirty read problem:

Transaction T 4 → A

bt := A.read();

A.write(bt+4);

A: 107

Transaction U 3 → A

bu := A.read();

A.write(bu+3);

Commit

Abort

Correct result: A = 103

T修改了数据但未提交，U读取了T修改过的数据；

如果T回滚，U读取的值会成为无效的脏数据

- premature write (over-writing uncommitted values)

- Over-writing uncommitted values :

Transaction T 4 → A

bt := A.read();

A.write(bt+4);

A: 107

Transaction U 3 → A

bu := A.read();

A.write(bu+3);

Abort

Correct result: A = 104

事务 U 回滚后，其修改无效，但它已经覆盖了事务 T 的未提交值，导致事务 T 的更新丢失；

- solutions

- **cascading aborts:** a transaction reading uncommitted data must be aborted if the transaction that modified the data aborts
- to avoid cascading aborts (too many rollbacks)  
transactions can only read data written by committed transactions
- how to preserve data despite subsequent failures  
stable storage  
2 copies of data stored in separate parts of disks && not decay related  
decay related: probability of both parts corrupted is small

## nested transactions

- transactions composed of **sub-transactions**
- subtransactions commit / abort independently
- effect of sub-transaction becomes durable only when top-level transaction commits

## concurrency control: locking

- environment:  
shared data in a single server, competing clients
- goal:  
transactions && maximizing concurrency
- solution:
- serial equivalence

## exclusive locks

- Exclusive locks

**Transaction T** A → B: 4

bt := A.read();

A.write(bt-4);

bt := B.read();

B.write(bt+4);

A: 96

B: 200

C: 297

**Transaction U** C → B: 3

bu := C.read();

C.write(bu-3);

bu := B.read();

Wait for T

- Exclusive locks

**Transaction T** A → B: 4

bt := A.read();

A.write(bt-4);

bt := B.read();

B.write(bt+4);

CloseTransaction(T);

A: 96

B: 204

C: 297

**Transaction U** C → B: 3

bu := C.read();

C.write(bu-3);

bu := B.read();

Wait for T

- basic elements of protocol

1. serial equivalence -> 2-phase locking

- growing\_phase

加锁阶段, 拥有了锁之后可以继续加锁

确保transaction锁住所有需要的资源

- shrinking\_phase

一旦transaction开始释放任何锁, 就进入了释放锁阶段, 不能再加锁

- lock compatability for 2-phase locking

lock state of the target data	action
not locked yet	lock set & operation proceeds
conflicting lock set by another transaction	wait till release
non-conflicting lock set by another transaction	lock shared & operation proceeds
locked by itself	lock promoted if necessary (read -> write) & operation proceeds

## 2. hide intermediate results

problem:

如果一个transaction在commit / abort之前释放锁，其他transaction可能访问到其未完成的中间结果

solution: strict 2-phase locking

better release of locks only at commit / abort, which means that locks held till end of transaction

### ▪ how to increase concurrency & preserve serial equivalence

#### ▪ granularity of locks 锁的粒度 (数据范围)

large granularity -> limits concurrent access

small granularity -> higher managing overhead

#### ▪ appropriate locking rules

lock compatibility

For one data item		Lock requested	
		Read	Write
Lock already set	None	OK	OK
	Read	OK	Wait
	Write	Wait	Wait

- lock implementation
  - lock manager
  - managing table of locks  
(which contains transaction identifiers, data item identifiers, lock type [shared / exclusive], condition variable)
- deadlocks

## Concurrency control: locking

- Read/write lock

**Transaction T** A → B: 4

bt := A.read();  
A.write(bt-4);

bt := B.read();

Wait for release by U

B.write(bt+4);

**Transaction U** C → B: 3

bu := C.read();

C.write(bu-3);

bu := B.read();

Wait for release by T

B.write(bu+3);

**Deadlock!!**

- prevention
  - locking all data items used by a transaction when it starts
  - requesting locks on data items in a predefined order

but impossible for interactive transactions, (在interactive环境下用户会动态决定transaction的行为, 很难预先知道transaction将访问哪些数据) && reduction of concurrency

- detection

server keeps track of a wait-for graph,

lock - edge added,

unlock - edge removed



checks **cycles** when an edge is added

■ solution

1. once a deadlock detected, server selects a transaction and aborts it (**to break the cycle**)

and the choice? factors:

age of transaction, \_\_ (年龄较小的transaction重启成本较低) \_\_

number of cycles the transaction is involved in (优先终止涉及更多deadlock的 transaction) \_\_

2. timeouts

locks granted for a limited period of time

within period: lock invulnerable (**non-preemptive**)

after period: lock vulnerable (**preemptive**)

存在中断有效transaction的风险, 适合实时性要求高, transaction复杂度低的场景